

---

# Sculpting Vulnerability Basins: Physical Sharpness-Aware Minimization for Embodied VLAs

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Vision-Language-Action (VLA) models have emerged as powerful paradigms for embodied robotic control, yet their robustness against physical-world adversarial attacks remains underexplored. Current adversarial patch generation methods typically optimize over static, isolated 2D frames. However, deploying these patches against robotic systems exposes a critical flaw: standard optimizations tend to converge into fragile sharp minima. Consequently, these attacks suffer from severe adversarial flickering and fail completely when subjected to macroscopic kinematic shifts (e.g., continuous camera translation) and microscopic hardware degradations (e.g., printer gamut limits). To bridge this sim-to-real gap, we propose **Physical Sharpness-Aware Minimization (P-SAM)**. Our framework mathematically decouples physical robustness into two tractable objectives. First, we model macroscopic trajectory dynamics using transformation expectations over offline datasets. Second, we utilize the white-box nature of the VLA to perform a minimax optimization, dynamically approximating the conditionally worst-case microscopic pixel noise via first-order Taylor expansions. By forcing gradient updates at these worst-case perturbed states, P-SAM actively sculpts broad vulnerability basins in the loss landscape. Combined with strict physical fabricability constraints, our method generates patches that maintain inescapable, devastating influence over the VLA’s 7-DoF kinematics. Evaluations across diverse manipulation tasks demonstrate that P-SAM significantly outperforms baseline static patches, underscoring an urgent need for robust defensive paradigms in embodied AI. Our project can be found at <https://psam-26.github.io>.

## 1 Introduction

The integration of high-capacity vision encoders with large language models [Deitke et al., 2025, Beyer et al., 2024, Chen et al., 2023, Qwen et al., 2025, Karamcheti et al., 2024, Alayrac et al., 2022, Wang et al., 2024, Yao et al., 2024] has catalyzed a paradigm shift in embodied artificial intelligence. Vision-Language-Action (VLA) foundation models [Wang et al., 2025b, Kim et al., 2024, Black et al., 2026, Ding et al., 2025, Shi et al., 2025, Team et al., 2025] can directly map raw visual observations and natural language instructions to continuous, low-level robotic kinematics (e.g., 7-Degree-of-Freedom control). By fusing semantic reasoning with physical execution in an end-to-end manner, VLAs enable robots to perform complex, zero-shot manipulation tasks in unstructured, open-world environments. However, as these models transition from simulated benchmarks to safety-critical physical deployments [Liu et al., 2025, NVIDIA et al., 2025, Intelligence et al., 2025, Smith and Coles, 1973, Kroemer et al., 2020, Ingrand and Ghallab, 2017], their direct reliance on unmediated visual inputs introduces a profound attack surface. Understanding and mitigating their vulnerabilities against malicious physical-world adversarial interference has consequently emerged as a paramount security imperative.




Figure 1: Overcoming adversarial flickering in embodied VLAs. (a) Standard patches optimized on static frames fail in closed-loop systems, as macroscopic kinematic shifts and microscopic hardware noise instantly neutralize the attack. (b) While existing methods converge to fragile sharp minima (left), our P-SAM sculpts robust flat minima (right), maintaining attack efficacy despite unmodeled physical perturbations  $\epsilon$ . (c) TV and NPS constraints ensure physical fabricability, preventing color clipping and guaranteeing real-world deployment.

Historically, physical adversarial attacks (e.g., printable patches placed in the environment) have proven highly effective at deceiving static computer vision models [Eykholt et al., 2018a,b, Wei et al., 2024, DelaCruz et al., 2026, Nguyen et al., 2024, Wang et al., 2023]. Yet, deploying these conventional patches against an embodied, closed-loop robotic system exposes a fundamental methodological flaw. Current adversarial patch generation methods [Wang et al., 2025a, Xu et al., 2020a, Huang et al., 2026] typically optimize perturbations over isolated, static 2D frames. However, a robotic manipulator operates as a continuous dynamical system. As the robot moves, the patch undergoes macroscopic kinematic shifts (e.g., continuous changes in camera translation, scale, and viewing angle) and is subjected to microscopic hardware degradations (e.g., CMOS sensor grain, motion blur, and printer color gamut clipping). Because standard empirical risk minimization naturally converges into highly localized, sharp minima in the highly non-convex adversarial loss landscape, even a microscopic deviation instantly ejects the patch from this fragile optimal state. We characterize this phenomenon as adversarial flickering (see Figure 1 (a)), a state where the malicious patch succeeds in one frame but shatters the moment the robot arm moves, allowing the VLA’s closed-loop feedback mechanism to safely auto-correct its trajectory and override the attack.

Bridging this adversarial sim-to-real gap poses a formidable technical bottleneck. Naive solutions might attempt to simulate the continuous temporal dynamics of the robot using heavy 3D differentiable physics simulators or by unrolling full kinematic trajectories during optimization. However, such approaches introduce prohibitive computational overhead, complex sim-to-real dynamics modeling, and severe memory bottlenecks, rendering them highly impractical for scaling across diverse robotic embodiments and environments.

To overcome these challenges without the computational explosion of 3D digital twins, we propose **Physical Sharpness-Aware Minimization (P-SAM)**. We hypothesize that for a physical attack to reliably hijack a continuous control loop, the optimization must actively sculpt a broad, robust vulnerability basin (i.e., a flat minimum, as illustrated in Figure 1 (b)) rather than a sharp valley. Our computationally lightweight framework mathematically decouples physical robustness into two tractable objectives entirely within the 2D domain. First, to account for macroscopic kinematic shifts, we utilize Expectation over Transformation (EoT) on single frames sampled from offline robotic datasets, cheaply simulating the continuous camera motion of the robot arm traversing 3D space. Second, to combat microscopic hardware noise, P-SAM formulates a minimax objective utilizing the white-box computational graph of the VLA. We dynamically approximate the conditionally worst-case pixel perturbation for any given macroscopic perspective via a first-order Taylor expansion, requiring only a single backward pass. By forcing the gradient descent step to evaluate at these worst-case perturbed states, our optimization acts as a mathematical repellent, dragging the patch parameters away from fragile cliffs and explicitly flattening the adversarial loss landscape.

To guarantee that this mathematically optimal flat minimum translates seamlessly to the physical world, we strictly constrain the optimization using Total Variation (TV) and Non-Printability Score

(NPS) regularizers. These constraints ensure the resulting patches feature organic, smooth color transitions that survive motion blur and are perfectly reproducible by standard commercial printers (see Figure 1 (c)). The final patches maintain an inescapable, devastating influence over the VLA’s closed-loop kinematics, bypassing the need for digital access to the target robot.

In summary, our core contributions are as follows:

- We expose a critical failure mode of current static adversarial patches against VLAs, mathematically attributing the physical sim-to-real gap to the convergence of standard optimizers into brittle sharp minima.
- We propose a novel, computationally lightweight framework that elegantly decouples macroscopic robotic kinematics (via offline 2D EoT) from microscopic hardware noise. By forcing gradient updates at conditionally worst-case perturbations, P-SAM actively sculpts broad vulnerability basins without requiring heavy 3D differentiable rendering.
- By integrating P-SAM with strict physical fabricability constraints, we generate adversarial patches that reliably and consistently hijack continuous 7-DoF robotic manipulation tasks in the real world. Extensive evaluations demonstrate that our method drastically outperforms existing baseline static patches, underscoring an urgent need for robust defensive architectures in next-generation VLA models.

## 2 Method

Current adversarial patch generation methods for Vision-Language-Action (VLA) models predominantly optimize perturbations over static, isolated 2D images. This static assumption fundamentally ignores the dynamic realities of embodied robotic systems. During real-world execution, an adversarial patch is continuously subjected to macroscopic kinematic shifts (e.g., motion distortion, scaling, and perspective changes as the camera moves) and microscopic hardware degradations (e.g., CMOS sensor noise, ambient lighting jitter, and printer color gamut clipping). By ignoring these physical factors, conventional static 2D optimizations converge into brittle solutions that suffer from adversarial flickering, a phenomenon where a patch succeeds in a pristine digital frame but fails the moment the robotic arm moves.

To bridge this critical sim-to-real gap, we propose a novel framework that systematically embeds physical-world dynamics into the 2D optimization pipeline (as shown in Figure 2). Our method mathematically decouples physical robustness into three tractable objectives: simulating macroscopic motion dynamics (Section 2.1), sculpting flat minima to absorb microscopic hardware noise (Section 2.2), and bounding the optimization with strict physical fabricability constraints to guarantee zero-shot real-world transferability (Section 2.3).

### 2.1 Macroscopic Kinematic Simulation via Expectation over Transformation

Conventional adversarial attacks generate patches by overfitting to pristine, static viewpoints. The moment the robotic arm initiates a trajectory, the resulting perspective distortion (changes in scale, viewing angle, and spatial warping) instantly breaks the adversarial alignment. To guarantee robustness against these macroscopic kinematic shifts without leaving the efficient 2D domain, we employ Expectation over Transformation (EoT) across offline workspace datasets (see Figure 2 (a)).

Formally, the physically augmented observation fed into the VLA is defined as:

$$o(x, \delta, \tau) = (1 - M_\tau) \odot x + M_\tau \odot \mathcal{T}(\delta, \tau), \quad (1)$$

where  $x \in \mathbb{R}^{H \times W \times 3}$  is a clean, benign camera frame sampled from an offline dataset (e.g., LIBERO), and  $\delta$  represents the learnable parameters of the 2D adversarial patch. To simulate the continuous perspective distortions and motion artifacts caused by the robot’s movement, we sample a physical transformation  $\tau$  from a predefined physical distribution  $\Gamma$  (encompassing random affine scaling, perspective warping, rotation, and brightness shifts). The function  $\mathcal{T}(\delta, \tau)$  applies this macroscopic transformation to the digital patch  $\delta$ , while the operator  $\odot$  denotes the Hadamard (element-wise) product.

Crucially, the dynamically transformed binary mask  $M_\tau \in \{0, 1\}^{H \times W}$  is not merely a spatial bounding box, but the mathematical guarantor of physical realizability. To distinguish a legitimate




Figure 2: Overview of the P-SAM framework. Our method mitigates the adversarial sim-to-real gap via three decoupled 2D objectives. (a) Macroscopic Kinematic Simulation: Expectation over Transformation (EoT) with dynamic opaque masking models continuous 3D perspective shifts. (b) Microscopic Flat-Minima Optimization: Evaluating gradients at worst-case perturbed states sculpts a flat vulnerability basin to mitigate hardware-induced adversarial flickering. (c) Physical Fabricability Constraints:  $\mathcal{L}_{TV}$  and  $\mathcal{L}_{NPS}$  regularizers ensure the patch survives motion blur and CMYK gamut clipping during physical deployment.

localized physical attack from an impossible global digital perturbation,  $M_\tau$  enforces two strict constraints: opaque occlusion and dynamic geometric synchronization. To prevent physically impossible semi-transparent overlays,  $(1 - M_\tau) \odot x$  acts as a spatial carving mechanism that precisely deletes background pixels before patch insertion. Simultaneously, applying the identical macroscopic transformation  $\tau$  to the base mask ensures that  $M_\tau$  tracks the patch’s perspective distortions. This prevents out-of-bound digital artifacts and strictly confines adversarial gradients to the valid physical dimensions of the deployable sticker. A comprehensive mathematical and geometric analysis of these constraints is provided in Appendix C.3.

Given this strictly bounded, physically augmented observation, we formulate the base targeted manipulation loss:

$$\mathcal{L}_{base}(\delta, x, \tau) = \mathcal{L}_{CE}(\mathcal{F}_{VLA}(o(x, \delta, \tau), I), A_{target}), \quad (2)$$

where  $\mathcal{F}_{VLA}(\cdot)$  acts as a frozen white-box VLA model, whose internal weights are fixed, but full access to its computational graph is maintained to backpropagate gradients.  $\mathcal{L}_{CE}$  denotes the cross-entropy loss, which quantifies the discrepancy between the VLA’s predicted action distribution (conditioned on the natural language instruction  $I$ ) and the attacker’s malicious 7-DoF target  $A_{target}$ .

## 2.2 Conquering Microscopic Noise via Physical Sharpness-Aware Minimization

While EoT addresses macroscopic viewpoint changes, standard empirical risk minimization over static datasets inevitably drives the patch parameters into a sharp minimum. In the physical world, cameras do not capture perfect pixel matrices. Unmodeled microscopic variations (e.g., CMOS camera sensor grain, ambient lighting fluctuations, and slight printer ink bleeding) act as random noise vectors that instantly eject the patch from this fragile minimum, neutralizing the attack.

To actively sculpt a globally robust vulnerability basin (i.e., a flat minimum) that immunizes the patch against these unpredictable sensor and environmental noises (as shown in Figure 2 (b)), we introduce Physical Sharpness-Aware Minimization (P-SAM). We enforce a constraint that the patch must remain aggressively malicious even when subjected to the absolute worst-case microscopic pixel degradation. This is formulated as a minimax objective:

$$\min_{\delta} \mathbb{E}_{x, \tau} \left[ \max_{\|\epsilon\|_2 \leq \rho} \mathcal{L}_{base}(\delta + \epsilon, x, \tau) \right], \quad (3)$$

where  $\epsilon$  is a microscopic, learnable perturbation applied strictly to the patch pixels, acting as a mathematical surrogate for unmodeled physical hardware noise. The hyperparameter  $\rho$  defines the  $\ell_2$ -norm radius of the perturbation ball, effectively controlling the desired width of the vulnerability basin.

Solving the exact inner maximization iteratively is computationally intractable for deep neural networks. However, utilizing the white-box gradients of the VLA, we can efficiently approximate

the conditionally worst-case noise  $\epsilon^*$  via a first-order Taylor expansion (see Appendix C.2 for the detailed mathematical derivation and physical justification), requiring only a single additional backward pass per iteration:

$$\epsilon^*(x, \tau) = \rho \frac{\nabla_{\delta} \mathcal{L}_{base}(\delta, x, \tau)}{\|\nabla_{\delta} \mathcal{L}_{base}(\delta, x, \tau)\|_2}, \quad (4)$$

where  $\nabla_{\delta} \mathcal{L}_{base}$  denotes the standard gradient of the base loss with respect to the patch parameters. Crucially, this worst-case pixel noise  $\epsilon^*(x, \tau)$  is strictly conditioned on the specific background image  $x$  and the sampled macroscopic perspective  $\tau$ .

Instead of updating the patch using the gradient at its current state  $\delta_k$ , P-SAM computes the gradient at this worst-case perturbed state. The patch parameters are thus updated as:

$$\delta_{k+1} = \delta_k - \alpha \cdot \mathbb{E}_{(x, \tau) \sim \mathcal{B}} [\nabla_{\delta} \mathcal{L}_{base}(\delta_k + \epsilon^*(x, \tau), x, \tau)], \quad (5)$$

where  $\mathcal{B}$  is the current mini-batch of sampled image-transformation pairs and  $\alpha$  is the optimization learning rate.

### Geometric Interpretation of Global Flatness:

As the local loss landscape varies continuously with respect to the macroscopic viewpoint  $\tau$ , the corresponding worst-case perturbation  $\epsilon^*$  exhibits significant variance. By independently computing the conditional worst-case perturbation for each  $(x, \tau)$  pair within a mini-batch and averaging their perturbed gradients, the optimization trajectory of the patch  $\delta$  simultaneously minimizes the loss across an ensemble of viewpoint-specific worst-case degradations. This dynamic effectively regularizes against sharp local minima, encouraging convergence toward a flatter, more resilient vulnerability basin (see Figure 3).




Figure 3: Comparison of flat and sharp minima under physical domain shifts. The domain shift is conceptually simulated by curve translation.

## 2.3 Physical Fabricability Constraints

Finally, a mathematically optimal digital patch is fundamentally useless if it cannot be accurately reproduced in the physical world. Static 2D optimization methods inherently assume perfect digital-to-physical pixel fidelity, frequently generating highly saturated, high-frequency noise patterns. Such patterns not only exceed the color gamut of physical printers (resulting in severe color clipping) but are also instantly destroyed by the motion blur inherent to a moving robotic camera. To bridge this final deployment gap, we enforce two strict physical regularizers (as shown in Figure 2 (c)).

First, to penalize high-frequency color transitions and ensure the adversarial signal survives motion blur and camera down-sampling, we apply Total Variation (TV) regularization:

$$\mathcal{L}_{TV}(\delta) = \sum_{i,j} \sqrt{(\delta_{i+1,j} - \delta_{i,j})^2 + (\delta_{i,j+1} - \delta_{i,j})^2}, \quad (6)$$

which guarantees organic, smooth gradients across the spatial coordinates  $(i, j)$  of the patch.

Second, to ensure the digital pixels can be faithfully reproduced by physical ink without distortion, we apply the Non-Printability Score (NPS):

$$\mathcal{L}_{NPS}(\delta) = \sum_{p \in \delta} \min_{c \in C} \|p - c\|_2, \quad (7)$$

which strictly binds every digital pixel vector  $p$  to the closest matching color  $c$  within a predefined, physically reproducible CMYK printer color gamut  $C$ . This prevents unpredictable out-of-gamut color clipping during fabrication.

**Final Unified Objective.** Integrating the macroscopic kinematic simulation (EoT), the microscopic flat-minima optimization (P-SAM), and the physical fabricability constraints, the complete, end-to-end objective function optimized by the attacker is:

$$\mathcal{L}_{final}(\delta) = \mathbb{E}_{(x, \tau) \sim \mathcal{B}} [\mathcal{L}_{base}(\delta + \epsilon^*(x, \tau), x, \tau)] + \lambda_1 \mathcal{L}_{TV}(\delta) + \lambda_2 \mathcal{L}_{NPS}(\delta), \quad (8)$$



---

**Algorithm 1** Physical Sharpness-Aware Minimization (P-SAM)

---

**Input:** Offline dataset  $\mathcal{D}$ , Frozen VLA  $\mathcal{F}_{VLA}$  with language instruction  $I$ , Target Action  $A_{target}$

**Parameters:** Learning rate  $\alpha$ , perturbation bound  $\rho$ , constraint weights  $\lambda_1, \lambda_2$

**Initialize:** Patch parameters  $\delta_0$

```
1: for step  $k = 0, 1, 2, \dots$  do
2:   Sample a mini-batch of clean frames  $x \sim \mathcal{D}$  and physical transforms  $\tau \sim \Gamma$ 
3:   Construct augmented observations  $o = (1 - M_\tau) \odot x + M_\tau \odot \mathcal{T}(\delta_k, \tau)$ 
4:   // Step 1: Inner Maximization (Approximate worst-case noise)
5:   Compute base loss  $\mathcal{L}_{base} = \mathcal{L}_{CE}(\mathcal{F}_{VLA}(o, I), A_{target})$ 
6:   Calculate worst-case perturbation:  $\epsilon^* = \rho \frac{\nabla_{\delta} \mathcal{L}_{base}}{\|\nabla_{\delta} \mathcal{L}_{base}\|_2}$ 
7:   // Step 2: Outer Minimization (Evaluate at perturbed state)
8:   Apply worst-case noise to the patch:  $\hat{\delta} = \delta_k + \epsilon^*$ 
9:   Construct noisy observation  $o_{noisy} = (1 - M_\tau) \odot x + M_\tau \odot \mathcal{T}(\hat{\delta}, \tau)$ 
10:  Compute perturbed loss  $\hat{\mathcal{L}}_{base} = \mathcal{L}_{CE}(\mathcal{F}_{VLA}(o_{noisy}, I), A_{target})$ 
11:  Compute gradients of constraints:  $\nabla_{\delta} \mathcal{L}_{TV}(\delta_k)$  and  $\nabla_{\delta} \mathcal{L}_{NPS}(\delta_k)$ 
12:  // Step 3: Apply Flat-Minima Update
13:  Compute final gradient:  $g = \nabla_{\delta} \hat{\mathcal{L}}_{base} + \lambda_1 \nabla_{\delta} \mathcal{L}_{TV}(\delta_k) + \lambda_2 \nabla_{\delta} \mathcal{L}_{NPS}(\delta_k)$ 
14:  Update patch:  $\delta_{k+1} = \delta_k - \alpha \cdot g$ 
15:  Project  $\delta_{k+1}$  to valid RGB color space  $[0, 1]$ 
16: end for
17: Return: Optimized, physically robust adversarial patch  $\delta^*$ 
```

---

where  $\lambda_1$  and  $\lambda_2$  are empirically tuned weighting coefficients that balance the targeted kinematic hijacking success rate against the stealth and physical printability of the adversarial patch. The full optimization procedure is summarized in Algorithm 1. By optimizing Equation 8 entirely offline, the attacker bypasses the need for computationally prohibitive 3D simulations or real-time queries to the target robot. The resulting P-SAM patch acts as a condensed vulnerability basin, inherently robust to both macroscopic kinematic shifts and microscopic physical degradations.

## 3 Experiments

### 3.1 Experimental Setup

**Victim Models and Datasets.** We evaluate our framework against OpenVLA-7B [Kim et al., 2024] and  $\pi_0$  [Black et al., 2026], prominent examples of open-weights Vision-Language-Action foundation models. LIBERO [Liu et al., 2023] is a simulation dataset designed to evaluate models across four distinct task types (Spatial, Object, Goal, Long). Notably, LIBERO-Spatial emphasizes spatial reasoning and geometric understanding, requiring models to comprehend relative positions and scene layouts, while LIBERO-Long combines diverse objects, layouts, and long-horizon tasks, making it challenging for complex, multi-step planning. Therefore, we utilize both the LIBERO-Long and LIBERO-Spatial suites to comprehensively assess these capabilities.

**Baselines and Objectives.** To systematically decompose the physical sim-to-real gap and validate the necessity of each component in our framework, we establish the following contrastive baselines and ablation configurations. All optimized patches utilize a targeted manipulation objective to maximize the continuous kinematic deviation from the benign trajectory: **(1) Clean (Initial Model):** Normal robotic execution without any adversarial interference, establishing the foundational capabilities of the VLA. **(2) Random Noise:** A patch instantiated with pixel values uniformly sampled from a Gaussian distribution, evaluating the VLA’s natural robustness against unstructured visual clutter. **(3) w/o Viewpoint Deformation:** A static adversarial patch optimized via standard Projected Gradient Descent (PGD) on isolated 2D frames, strictly ignoring continuous macroscopic kinematic shifts. **(4) w/o Flat Minima:** A patch that incorporates macroscopic Expectation over Transformation (EoT) but omits our microscopic minimax optimization, thus naturally converging to a fragile sharp minimum. **(5) w/o  $\mathcal{L}_{TV}$ :** The P-SAM framework excluding the Total Variation regularizer, permitting high-frequency spatial pixel transitions. **(6) w/o  $\mathcal{L}_{NPS}$ :** The P-SAM framework excluding the Non-Printability Score constraint, freely generating out-of-gamut digital colors. **(7) Ours (Full P-SAM):** Our complete, unified framework encompassing macroscopic EoT, microscopic worst-case minimax optimization, and strict physical fabricability constraints.

Table 1: Comparison of P-SAM and baselines. In Real-World scenarios, we employ (TFR  $\uparrow$ ) as the evaluation metric. OBJ, PLC, and MNP denote object grasping, placement, and manipulation.

| Method Configuration                    | Model   | LIBERO-Long |             |             | LIBERO-Spatial |             |             | Real-World |           |           |
|-----------------------------------------|---------|-------------|-------------|-------------|----------------|-------------|-------------|------------|-----------|-----------|
|                                         |         | TFR (%)     | NAD (%)     | ASR (%)     | TFR (%)        | NAD (%)     | ASR (%)     | OBJ (%)    | PLC (%)   | MNP (%)   |
| Clean (Initial Model)                   | Openvla | 46.3        | -           | -           | 15.3           | -           | -           | 12         | 14        | 20        |
| Random Noise                            |         | 48.4        | -           | -           | 16.6           | -           | -           | 14         | 17        | 21        |
| w/o Viewpoint Deformation               |         | 60.6        | 30.8        | 19.5        | 38.2           | 16.8        | 8.9         | 37         | 43        | 48        |
| w/o Flat Minima                         |         | 86.8        | 42.7        | 29.0        | 71.6           | 26.1        | 18.3        | 40         | 49        | 45        |
| w/o $\mathcal{L}_{TV}$ (No Smoothness)  |         | 72.4        | 34.9        | 25.4        | 49.2           | 14.5        | 10.7        | 35         | 44        | 47        |
| w/o $\mathcal{L}_{NPS}$ (No Gamut Sync) |         | 98.2        | 61.5        | 44.6        | 88.4           | 45.3        | 36.9        | 26         | 28        | 29        |
| <b>Ours (Full P-SAM)</b>                |         | <b>100</b>  | <b>63.6</b> | <b>47.1</b> | <b>95.6</b>    | <b>50.7</b> | <b>38.5</b> | <b>62</b>  | <b>71</b> | <b>76</b> |
| Clean (Initial Model)                   | $\pi_0$ | 6.2         | -           | -           | 1.0            | -           | -           | 1          | 4         | 7         |
| Random Noise                            |         | 8.2         | -           | -           | 1.8            | -           | -           | 2          | 4         | 9         |
| w/o Viewpoint Deformation               |         | 31.4        | 12.5        | 6.2         | 18.8           | 8.3         | 4.5         | 16         | 20        | 26        |
| w/o Flat Minima                         |         | 37.8        | 16.4        | 9.1         | 35.2           | 15.9        | 8.7         | 29         | 28        | 30        |
| w/o $\mathcal{L}_{TV}$ (No Smoothness)  |         | 36.0        | 16.1        | 9.0         | 36.2           | 16.2        | 8.8         | 27         | 29        | 34        |
| w/o $\mathcal{L}_{NPS}$ (No Gamut Sync) |         | 59.0        | 25.7        | 16.8        | 54.4           | 23.6        | 16.4        | 11         | 17        | 22        |
| <b>Ours (Full P-SAM)</b>                |         | <b>72.6</b> | <b>28.9</b> | <b>20.5</b> | <b>71.6</b>    | <b>29.3</b> | <b>20.2</b> | <b>43</b>  | <b>45</b> | <b>58</b> |

**Evaluation Metrics.** Unlike static image classification, evaluating robotic hijacking requires continuous temporal metrics: **(1) Task Failure Rate (TFR  $\uparrow$ ):** The percentage of task rollouts that fail to complete the designated instruction. **(2) Normalized Action Deviation (NAD  $\uparrow$ ):** The continuous  $\ell_2$ -norm distance between the adversarially predicted 7-DoF actions and the unperturbed ground-truth trajectory, normalized by the empirical operational range of each DoF. This strictly quantifies the true severity of the ongoing spatial and rotational hijacking. **(3) Attack Sustaining Rate (ASR  $\uparrow$ ):** The percentage of sequential frames within a rollout where the adversarial patch successfully induces an NAD greater than a critical safety threshold ( $\beta = 0.1$ ). A low ASR indicates the attack repeatedly loses its grip, allowing the VLA to auto-correct and return to its benign path. See Appendix A and Appendix E for more details.

### 3.2 Main Result

Table 1 summarizes quantitative evaluations across continuous noisy simulations and real-world physical deployments. While **Clean** models and **Random Noise** confirm the inherent resilience of VLAs against unstructured visual anomalies, conventional adversarial patches expose a profound sim-to-real gap. For instance, the static **w/o Viewpoint Deformation** baseline suffers a catastrophic real-world collapse (e.g., ASR plunges to 19.5% on OpenVLA and 6.2% on  $\pi_0$ ), empirically capturing adversarial flickering: static attacks instantly misalign during continuous camera translation. Although the **w/o Flat Minima** patch survives macroscopic shifts, its sharp optimum remains brittle to unmodeled hardware noise, allowing the VLA to auto-correct and stalling real-world TFRs at 40%–49% for OpenVLA. Furthermore, digital robustness is useless without fabrication constraints. The **w/o  $\mathcal{L}_{NPS}$**  baseline exhibits a textbook sim-to-real collapse (98.2% digital to 26%–29% physical TFR) due to irreversible CMYK color clipping, while the **w/o  $\mathcal{L}_{TV}$**  patch generates high-frequency static that is easily destroyed by natural motion blur acting as a low-pass filter.

In stark contrast, **Ours (Full P-SAM)** acts as a mathematical shock absorber. By actively sculpting a physically bounded flat vulnerability basin, P-SAM sustains a 100% digital TFR on OpenVLA-Long and maximizes real-world TFRs up to 76% in complex Manipulation (MNP) tasks. Even against the resilient  $\pi_0$  architecture, P-SAM successfully scales the physical failure rate to 58%. Crucially, P-SAM nearly doubles the ASR compared to baseline methods (e.g., reaching 47.1% on OpenVLA-Long), ensuring that the continuous adversarial hijacking remains inescapable throughout the physical execution.

### 3.3 More Analysis

**Spatial Sensitivity Tests.** A critical dimension of real-world threat modeling is whether an adversarial patch requires strict physical proximity to the target object to succeed. Taking OpenVLA as the target model, we systematically varied the patch placement across nine distinct angular directions (Figure 4 (a)) and incrementally increased its distance offset up to 30 centimeters (Figure 4 (b)). Strikingly, P-SAM maintains a consistently high TFR across all angular zones, experiencing negligible degradation even when pushed to the peripheral boundaries of the workspace. Because modern VLA vision encoders rely heavily on global self-attention mechanisms to construct holistic scene representations, the optimized patch effectively acts as an “attention sink” that corrupts




Figure 4: TFR across different spatial patch placements. (a) TFR across diverse angular placements. (b) TFR under increasing distance offsets.

the global visual feature space from any visible location. This spatial and distance invariance underscores the stealth of the P-SAM attack, explicitly demonstrating that adversaries can deploy the patch in highly inconspicuous peripheral regions without compromising continuous attack efficacy.

**Hyperparameter Evaluations.** The perturbation radius  $\rho$  in our minimax objective dictates the width of the sculpted vulnerability basin. As shown in Figure 5 (a), increasing  $\rho$  from 0.01 to 0.05 significantly raises the real-world TFR (e.g., scaling up to 76% in OpenVLA manipulation tasks and 71% in OpenVLA placement tasks), demonstrating enhanced robustness against noise. However, overly aggressive bounds ( $\rho > 0.1$ ) excessively restrict the optimization space, leading to underfitting and a subsequent decline in overall targeted efficacy. We empirically select  $\rho = 0.05$  as the optimal trade-off between physical resilience and adversarial severity. We also investigate the relationship between the adversarial patch size and hijacking efficacy in Figure 5 (b). We observe a direct correlation: as the patch size increases from 2% to 10% of the observation frame, the TFR rises across both models. While larger patches provide a broader canvas for the optimizer to embed robust spatial features, we constrain our default patch size to 5% to strike a balance between environmental stealth and critical system disruption. More supplementary experiments can be found in Appendix D.




Figure 5: The TFR results of  $\rho$  and patch size. The result of the patch size is the average value of OBJ, PLC and MNP.

**Transferability Assessments.** In realistic security threats, attackers rarely possess perfect white-box access to the deployed VLA’s internal parameters. Optimization theory suggests that converging to flat minima mitigates overfitting to the surrogate model’s specific decision boundaries, thereby inherently enhancing adversarial transferability. To rigorously evaluate this black-box threat potential, we conducted extensive zero-shot evaluations across both distinct architectures and diverse environments. Specifically, we optimized the adversarial patch on a single model and deployed it against other models, including OpenVLA [Kim et al., 2024],  $\pi_0$  [Black et al., 2026],  $\pi_{0.5}$  [Intelligence et al., 2025], and FAST [Pertsch et al., 2025] (see Figure 6 (a)). Furthermore, we investigated cross-scenario transferability by generating the patch within a specific setting and assessing its attack efficacy across distinct LIBERO [Liu et al., 2023] task suites (Spatial, Object, Goal, and Long) (see Figure 6 (b)). We can find that our P-SAM maintained zero-shot physical transferability. By sculpting a wide vulnerability basin, P-SAM captures generalized semantic flaws in the underlying vision encoders, posing a realistic physical threat even under strict black-box and cross-domain assumptions.




Figure 6: The TFR results of transferability across different VLAs and tasks. (a) Testing on LIBERO-Spatial. (b) Testing on OpenVLA.

**Qualitative Demonstrations.** To intuitively illustrate the behavioral shifts induced by different patches, Figure 7 and Figure 8 provides visualizations of continuous execution. **(1) Simulation Instance:** In a “Pick up the black bowl in the top drawer of the wooden cabinet and place it on the plate” task, the conventional adversarial attack baseline causes an initial adversarial deviation. However, as the camera moves, perspective distortion and simulated noise trigger adversarial flick-






Figure 7: Qualitative results of the simulated scenario. The deployed model is OpenVLA.




Figure 8: Qualitative results of the real-world scenario. The deployed model is OpenVLA.

ering; the VLA regains perception, auto-corrects the trajectory, and successfully completes the task. Conversely, P-SAM maintains a relentless adversarial grip, forcing the arm into a persistent erratic trajectory. (2) **Real-World Instance:** Deployed on the PiPER arm for a “Place the blue square on top of the red one” task, the conventional adversarial attack baseline patch fails within the first few seconds due to ambient lighting shifts and motion blur. However, the P-SAM patch effectively overrides the intended semantic instruction and subverts the 7-DoF closed-loop kinematics, thereby exposing critical physical security vulnerabilities. Appendix G presents more results.

## 4 Related Work

Adversarial attacks serve as indispensable tools for rigorously evaluating model vulnerabilities in the dynamic physical environments where robotic systems operate. While conventional gradient-driven pixel-level perturbations [Wang et al., 2019, Su et al., 2019, Madry et al., 2019, Liu et al., 2019, Goodfellow et al., 2015, Du et al., 2021] achieve high efficacy in static digital simulations, they fundamentally lack physical viability. Consequently, patch-based adversarial techniques [Wu et al., 2025, Xu et al., 2020b, Chen et al., 2023, Brown et al., 2018, Athalye et al., 2018, Billard and Kragic, 2019] have emerged as a more practical alternative by generating localized and physically fabricable perturbations. These patches remain robust against environmental shifts like changing illumination and camera angles. Given that Vision-Language-Action (VLA) models are designed for real-world execution, assessing their operational security through physical patches is highly critical.

## 5 Conclusion

In this work, we identified adversarial flickering as a primary cause of physical attack failures against Vision-Language-Action (VLA) models, demonstrating that static patches inherently converge into fragile, sharp minima. To address this, we proposed P-SAM, a framework that actively sculpts flat vulnerability basins by coupling macroscopic 2D perspective augmentations with a conditionally worst-case microscopic minimax optimization. Enforced by strict fabricability constraints, P-SAM generates deployable patches capable of reliably hijacking continuous robotic manipulation tasks in the physical world. By exposing critical vulnerabilities in VLA models that caution against impetuous real-world deployments, our framework serves as a pioneering contribution toward the reliability of AI-enhanced robotic systems.

**Limitations:** First, our offline patch generation process strictly requires white-box gradient access to the target VLA or an aligned surrogate to compute the worst-case pixel perturbations. Second, the physical attack inherently relies on continuous camera line-of-sight, meaning severe topological occlusions will temporarily suspend the adversarial influence.

## References

- Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katie Millican, Malcolm Reynolds, Roman Ring, Eliza Rutherford, Serkan Cabi, Tengda Han, Zhitao Gong, Sina Samangooei, Marianne Monteiro, Jacob Menick, Sebastian Borgeaud, Andrew Brock, Aida Nematzadeh, Sahand Sharifzadeh, Mikolaj Binkowski, Ricardo Barreira, Oriol Vinyals, Andrew Zisserman, and Karen Simonyan. Flamingo: a visual language model for few-shot learning, 2022. URL <https://arxiv.org/abs/2204.14198>.
- Anish Athalye, Logan Engstrom, Andrew Ilyas, and Kevin Kwok. Synthesizing robust adversarial examples, 2018. URL <https://arxiv.org/abs/1707.07397>.
- Lucas Beyer, Andreas Steiner, André Susano Pinto, Alexander Kolesnikov, Xiao Wang, Daniel Salz, Maxim Neumann, Ibrahim Alabdulmohsin, Michael Tschannen, Emanuele Bugliarello, Thomas Unterthiner, Daniel Keysers, Skanda Koppula, Fangyu Liu, Adam Grycner, Alexey Gritsenko, Neil Houlsby, Manoj Kumar, Keran Rong, Julian Eisenschlos, Rishabh Kabra, Matthias Bauer, Matko Bošnjak, Xi Chen, Matthias Minderer, Paul Voigtlaender, Ioana Bica, Ivana Balazevic, Joan Puigcerver, Pinelopi Papalampidi, Olivier Henaff, Xi Xiong, Radu Soricut, Jeremiah Harmsen, and Xiaohua Zhai. Paligemma: A versatile 3b vlm for transfer, 2024. URL <https://arxiv.org/abs/2407.07726>.
- Aude Billard and Danica Kragic. Trends and challenges in robot manipulation. *Science*, 364(6446): eaat8414, 2019. doi: 10.1126/science.aat8414. URL <https://www.science.org/doi/abs/10.1126/science.aat8414>.
- Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Lucy Xiaoyang Shi, James Tanner, Quan Vuong, Anna Walling, Haohuan Wang, and Ury Zhilinsky.  $\pi_0$ : A vision-language-action flow model for general robot control, 2026. URL <https://arxiv.org/abs/2410.24164>.
- Tom B. Brown, Dandelion Mané, Aurko Roy, Martín Abadi, and Justin Gilmer. Adversarial patch, 2018. URL <https://arxiv.org/abs/1712.09665>.
- Xi Chen, Xiao Wang, Lucas Beyer, Alexander Kolesnikov, Jialin Wu, Paul Voigtlaender, Basil Mustafa, Sebastian Goodman, Ibrahim Alabdulmohsin, Piotr Padlewski, Daniel Salz, Xi Xiong, Daniel Vlasic, Filip Pavetic, Keran Rong, Tianli Yu, Daniel Keysers, Xiaohua Zhai, and Radu Soricut. Pali-3 vision language models: Smaller, faster, stronger, 2023. URL <https://arxiv.org/abs/2310.09199>.
- Matt Deitke, Christopher Clark, Sangho Lee, Rohun Tripathi, Yue Yang, Jae Sung Park, Mohammadreza Salehi, Niklas Muennighoff, Kyle Lo, Luca Soldaini, Jiasen Lu, Taira Anderson, Erin Bransom, Kiana Ehsani, Huong Ngo, YenSung Chen, Ajay Patel, Mark Yatskar, Chris Callison-Burch, Andrew Head, Rose Hendrix, Favyen Bastani, Eli VanderBilt, Nathan Lambert, Yvonne Chou, Arnavi Chheda, Jenna Sparks, Sam Skjonsberg, Michael Schmitz, Aaron Sarnat, Byron Bischoff, Pete Walsh, Chris Newell, Piper Wolters, Tanmay Gupta, Kuo-Hao Zeng, Jon Borchardt, Dirk Groeneveld, Crystal Nam, Sophie Lebrecht, Caitlin Wittliff, Carissa Schoenick, Oscar Michel, Ranjay Krishna, Luca Weihs, Noah A. Smith, Hannaneh Hajishirzi, Ross Girshick, Ali Farhadi, and Aniruddha Kembhavi. Molmo and pixmo: Open weights and open data for state-of-the-art vision-language models. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 91–104, 2025. doi: 10.1109/CVPR52734.2025.00018.
- Miguel A. DelaCruz, Patricia Mae Santos, and Rafael T. Navarro. Physical adversarial attacks on ai surveillance systems: detection, tracking, and visible–infrared evasion, 2026. URL <https://arxiv.org/abs/2604.06865>.
- Pengxiang Ding, Han Zhao, Wenjie Zhang, Wenxuan Song, Min Zhang, Siteng Huang, Ningxi Yang, and Donglin Wang. Quar-vla: Vision-language-action model for quadruped robots, 2025. URL <https://arxiv.org/abs/2312.14457>.

395 Andrew Du, Bo Chen, Tat-Jun Chin, Yee Wei Law, Michele Sasdelli, Ramesh Rajasegaran, and  
396 Dillon Campbell. Physical adversarial attacks on an aerial imagery object detector, 2021. URL  
397 <https://arxiv.org/abs/2108.11765>.

398 Kevin Eykholt, Ivan Evtimov, Earlene Fernandes, Bo Li, Amir Rahmati, Florian Tramer, Atul  
399 Prakash, Tadayoshi Kohno, and Dawn Song. Physical adversarial examples for object detectors,  
400 2018a. URL <https://arxiv.org/abs/1807.07769>.

401 Kevin Eykholt, Ivan Evtimov, Earlene Fernandes, Bo Li, Amir Rahmati, Chaowei Xiao, Atul  
402 Prakash, Tadayoshi Kohno, and Dawn Song. Robust physical-world attacks on deep learning  
403 visual classification. In *2018 IEEE/CVF Conference on Computer Vision and Pattern Recogni-*  
404 *tion*, pages 1625–1634, 2018b. doi: 10.1109/CVPR.2018.00175.

405 Pierre Foret, Ariel Kleiner, Hossein Mobahi, and Behnam Neyshabur. Sharpness-aware minimiza-  
406 tion for efficiently improving generalization, 2021. URL [https://arxiv.org/abs/2010.](https://arxiv.org/abs/2010.01412)  
407 [01412](https://arxiv.org/abs/2010.01412).

408 Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial  
409 examples, 2015. URL <https://arxiv.org/abs/1412.6572>.

410 Yuanhao Huang, Yilong Ren, Jinlei Wang, Lujia Huo, Xuesong Bai, Jinchuan Zhang, and Haiyang  
411 Yu. Advreal: Physical adversarial patch generation framework for security evaluation of object  
412 detection systems. *Expert Systems with Applications*, 296:128967, 2026. ISSN 0957-4174. doi:  
413 10.1016/j.eswa.2025.128967. URL <http://dx.doi.org/10.1016/j.eswa.2025.128967>.

414 Félix Ingrand and Malik Ghallab. Deliberation for autonomous robots: A survey. *Ar-*  
415 *tificial Intelligence*, 247:10–44, 2017. ISSN 0004-3702. doi: [https://doi.org/10.1016/](https://doi.org/10.1016/j.artint.2014.11.003)  
416 [j.artint.2014.11.003](https://doi.org/10.1016/j.artint.2014.11.003). URL [https://www.sciencedirect.com/science/article/pii/](https://www.sciencedirect.com/science/article/pii/S0004370214001350)  
417 [S0004370214001350](https://www.sciencedirect.com/science/article/pii/S0004370214001350). Special Issue on AI and Robotics.

418 Physical Intelligence, Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess,  
419 Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Manuel Y. Galliker, Dibya Ghosh,  
420 Lachy Groom, Karol Hausman, Brian Ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Devin  
421 LeBlanc, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Allen Z.  
422 Ren, Lucy Xiaoyang Shi, Laura Smith, Jost Tobias Springenberg, Kyle Stachowicz, James Tanner,  
423 Quan Vuong, Homer Walke, Anna Walling, Haohuan Wang, Lili Yu, and Ury Zhilinsky.  $\pi_{0.5}$ : a  
424 vision-language-action model with open-world generalization, 2025. URL [https://arxiv.](https://arxiv.org/abs/2504.16054)  
425 [org/abs/2504.16054](https://arxiv.org/abs/2504.16054).

426 Siddharth Karamcheti, Suraj Nair, Ashwin Balakrishna, Percy Liang, Thomas Kollar, and Dorsa  
427 Sadigh. Prismatic vlms: Investigating the design space of visually-conditioned language models,  
428 2024. URL <https://arxiv.org/abs/2402.07865>.

429 Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair,  
430 Rafael Rafailov, Ethan Foster, Grace Lam, Pannag Sanketi, Quan Vuong, Thomas Kollar, Ben-  
431 jamin Burchfiel, Russ Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang, and Chelsea Finn.  
432 Openvla: An open-source vision-language-action model, 2024. URL [https://arxiv.org/](https://arxiv.org/abs/2406.09246)  
433 [abs/2406.09246](https://arxiv.org/abs/2406.09246).

434 Oliver Kroemer, Scott Niekum, and George Konidaris. A review of robot learning for manipu-  
435 lation: Challenges, representations, and algorithms, 2020. URL [https://arxiv.org/abs/](https://arxiv.org/abs/1907.03146)  
436 [1907.03146](https://arxiv.org/abs/1907.03146).

437 Aishan Liu, Xianglong Liu, Jiaxin Fan, Yuqing Ma, Anlan Zhang, Huiyuan Xie, and Dacheng Tao.  
438 Perceptual-sensitive gan for generating adversarial patches. In *Proceedings of the Thirty-Third*  
439 *AAAI Conference on Artificial Intelligence and Thirty-First Innovative Applications of Artificial*  
440 *Intelligence Conference and Ninth AAAI Symposium on Educational Advances in Artificial In-*  
441 *telligence*, AAAI’19/IAAI’19/EAAI’19. AAAI Press, 2019. ISBN 978-1-57735-809-1. doi:  
442 10.1609/aaai.v33i01.33011028. URL <https://doi.org/10.1609/aaai.v33i01.33011028>.

443 Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. Libero:  
444 Benchmarking knowledge transfer for lifelong robot learning. *arXiv preprint arXiv:2306.03310*,  
445 2023.

- 446 Jiaming Liu, Hao Chen, Pengju An, Zhuoyang Liu, Renrui Zhang, Chenyang Gu, Xiaoqi Li, Ziyu  
447 Guo, Sixiang Chen, Mengzhen Liu, Chengkai Hou, Mengdi Zhao, KC alex Zhou, Pheng-Ann  
448 Heng, and Shanghang Zhang. Hybridvla: Collaborative diffusion and autoregression in a unified  
449 vision-language-action model, 2025. URL <https://arxiv.org/abs/2503.10631>.
- 450 Hui Lu, Yi Yu, Yiming Yang, Chenyu Yi, Qixin Zhang, Bingquan Shen, Alex C. Kot, and Xudong  
451 Jiang. When robots obey the patch: Universal transferable patch attacks on vision-language-action  
452 models, 2026. URL <https://arxiv.org/abs/2511.21192>.
- 453 Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu.  
454 Towards deep learning models resistant to adversarial attacks, 2019. URL <https://arxiv.org/abs/1706.06083>.
- 456 Kien Nguyen, Tharindu Fernando, Clinton Fookes, and Sridha Sridharan. Physical adversarial at-  
457 tacks for surveillance: A survey. *IEEE Transactions on Neural Networks and Learning Systems*,  
458 35(12):17036–17056, 2024. doi: 10.1109/TNNLS.2023.3321432.
- 459 NVIDIA, :, Johan Bjorck, Fernando Castañeda, Nikita Cherniadev, Xingye Da, Runyu Ding,  
460 Linxi "Jim" Fan, Yu Fang, Dieter Fox, Fengyuan Hu, Spencer Huang, Joel Jang, Zhenyu Jiang,  
461 Jan Kautz, Kaushil Kundalia, Lawrence Lao, Zhiqi Li, Zongyu Lin, Kevin Lin, Guilin Liu,  
462 Edith Llonet, Loic Magne, Ajay Mandlekar, Avnish Narayan, Soroush Nasiriany, Scott Reed,  
463 You Liang Tan, Guanzhi Wang, Zu Wang, Jing Wang, Qi Wang, Jiannan Xiang, Yuqi Xie, Yinchen  
464 Xu, Zhenjia Xu, Seonghyeon Ye, Zhiding Yu, Ao Zhang, Hao Zhang, Yizhou Zhao, Ruijie Zheng,  
465 and Yuke Zhu. Gr00t n1: An open foundation model for generalist humanoid robots, 2025. URL  
466 <https://arxiv.org/abs/2503.14734>.
- 467 Karl Pertsch, Kyle Stachowicz, Brian Ichter, Danny Driess, Suraj Nair, Quan Vuong, Oier Mees,  
468 Chelsea Finn, and Sergey Levine. Fast: Efficient action tokenization for vision-language-action  
469 models, 2025. URL <https://arxiv.org/abs/2501.09747>.
- 470 Qwen, :, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan  
471 Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang,  
472 Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin  
473 Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li,  
474 Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang,  
475 Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report, 2025.  
476 URL <https://arxiv.org/abs/2412.15115>.
- 477 Lucy Xiaoyang Shi, Brian Ichter, Michael Equi, Liyiming Ke, Karl Pertsch, Quan Vuong, James  
478 Tanner, Anna Walling, Haohuan Wang, Niccolo Fusai, Adrian Li-Bell, Danny Driess, Lachy  
479 Groom, Sergey Levine, and Chelsea Finn. Hi robot: Open-ended instruction following with hier-  
480 archical vision-language-action models, 2025. URL <https://arxiv.org/abs/2502.19417>.
- 481 Michael H. Smith and L. Stephen Coles. Design of a low cost, general purpose robot. In *Proceedings*  
482 *of the 3rd International Joint Conference on Artificial Intelligence, IJCAI'73*, page 324–335, San  
483 Francisco, CA, USA, 1973. Morgan Kaufmann Publishers Inc.
- 484 Jiawei Su, Danilo Vasconcellos Vargas, and Kouichi Sakurai. One pixel attack for fooling deep  
485 neural networks. *IEEE Transactions on Evolutionary Computation*, 23(5):828–841, 2019. doi:  
486 10.1109/TEVC.2019.2890858.
- 487 Gemini Robotics Team, Saminda Abeyruwan, Joshua Ainslie, Jean-Baptiste Alayrac, Montser-  
488 rat Gonzalez Arenas, Travis Armstrong, Ashwin Balakrishna, Robert Baruch, Maria Bauza,  
489 Michiel Blokzijl, Steven Bohez, Konstantinos Bousmalis, Anthony Brohan, Thomas Buschmann,  
490 Arunkumar Byravan, Serkan Cabi, Ken Caluwaerts, Federico Casarini, Oscar Chang, Jose En-  
491 rique Chen, Xi Chen, Hao-Tien Lewis Chiang, Krzysztof Choromanski, David D’Ambrosio,  
492 Sudeep Dasari, Todor Davchev, Coline Devin, Norman Di Palo, Tianli Ding, Adil Dostmohamed,  
493 Danny Driess, Yilun Du, Debidatta Dwibedi, Michael Elabd, Claudio Fantacci, Cody Fong, Erik  
494 Frey, Chuyuan Fu, Marissa Giustina, Keerthana Gopalakrishnan, Laura Graesser, Leonard Hasen-  
495 clever, Nicolas Heess, Brandon Harnaez, Alexander Herzog, R. Alex Hofer, Jan Humprik, Atil  
496 Iscen, Mithun George Jacob, Deepali Jain, Ryan Julian, Dmitry Kalashnikov, M. Emre Karago-  
497 zler, Stefani Karp, Chase Kew, Jerad Kirkland, Sean Kirmani, Yuheng Kuang, Thomas Lampe,



Antoine Laurens, Isabel Leal, Alex X. Lee, Tsang-Wei Edward Lee, Jacky Liang, Yixin Lin, Sharath Maddineni, Anirudha Majumdar, Assaf Hurwitz Michaely, Robert Moreno, Michael Neunert, Francesco Nori, Carolina Parada, Emilio Parisotto, Peter Pastor, Acorn Pooley, Kanishka Rao, Krista Reymann, Dorsa Sadigh, Stefano Saliceti, Pannag Sanketi, Pierre Sermanet, Dhruv Shah, Mohit Sharma, Kathryn Shea, Charles Shu, Vikas Sindhwani, Sumeet Singh, Radu Soricut, Jost Tobias Springenberg, Rachel Sterneck, Razvan Surdulescu, Jie Tan, Jonathan Tompson, Vincent Vanhoucke, Jake Varley, Grace Vesom, Giulia Vezzani, Oriol Vinyals, Ayzaan Wahid, Stefan Welker, Paul Wohlhart, Fei Xia, Ted Xiao, Annie Xie, Jinyu Xie, Peng Xu, Sichun Xu, Ying Xu, Zhuo Xu, Yuxiang Yang, Rui Yao, Sergey Yaroshenko, Wenhao Yu, Wentao Yuan, Jingwei Zhang, Tingnan Zhang, Allan Zhou, and Yuxiang Zhou. Gemini robotics: Bringing ai into the physical world, 2025. URL <https://arxiv.org/abs/2503.20020>.

Donghua Wang, Wen Yao, Tingsong Jiang, Chao Li, and Xiaoqian Chen. Rfla: A stealthy reflected light adversarial attack in the physical world. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4432–4442, 2023. doi: 10.1109/ICCV51070.2023.00411.

Taowen Wang, Cheng Han, James Liang, Wenhao Yang, Dongfang Liu, Luna Xinyu Zhang, Qifan Wang, Jiebo Luo, and Ruixiang Tang. Exploring the adversarial vulnerabilities of vision-language-action models in robotics. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 6948–6958, October 2025a.

Weiyun Wang, Zhe Chen, Wenhao Wang, Yue Cao, Yangzhou Liu, Zhangwei Gao, Jinguo Zhu, Xizhou Zhu, Lewei Lu, Yu Qiao, and Jifeng Dai. Enhancing the reasoning ability of multimodal large language models via mixed preference optimization. *arXiv preprint arXiv:2411.10442*, 2024.

Zhibo Wang, Siyan Zheng, Mengkai Song, Qian Wang, Alireza Rahimpour, and Hairong Qi. advpattern: Physical-world attacks on deep person re-identification via adversarially transformable patterns. In *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 8340–8349, 2019. doi: 10.1109/ICCV.2019.00843.

Zhijie Wang, Zhehua Zhou, Jiayang Song, Yuheng Huang, Zhan Shu, and Lei Ma. Vlatest: Testing and evaluating vision-language-action models for robotic manipulation. *Proceedings of the ACM on Software Engineering*, 2(FSE):1615–1638, 2025b. ISSN 2994-970X. doi: 10.1145/3729343. URL <http://dx.doi.org/10.1145/3729343>.

Hui Wei, Hao Tang, Xuemei Jia, Zhixiang Wang, Hanxun Yu, Zhubo Li, Shin’ichi Satoh, Luc Van Gool, and Zheng Wang. Physical adversarial attack meets computer vision: A decade survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12):9797–9817, December 2024. ISSN 1939-3539. doi: 10.1109/tpami.2024.3430860. URL <http://dx.doi.org/10.1109/TPAMI.2024.3430860>.

Chen Henry Wu, Rishi Shah, Jing Yu Koh, Ruslan Salakhutdinov, Daniel Fried, and Aditi Raghunathan. Dissecting adversarial robustness of multimodal lm agents, 2025. URL <https://arxiv.org/abs/2406.12814>.

Kaidi Xu, Gaoyuan Zhang, Sijia Liu, Quanfu Fan, Mengshu Sun, Hongge Chen, Pin-Yu Chen, Yanzhi Wang, and Xue Lin. Adversarial t-shirt! evading person detectors in a physical world. In Andrea Vedaldi, Horst Bischof, Thomas Brox, and Jan-Michael Frahm, editors, *Computer Vision – ECCV 2020*, pages 665–681, Cham, 2020a. Springer International Publishing.

Kaidi Xu, Gaoyuan Zhang, Sijia Liu, Quanfu Fan, Mengshu Sun, Hongge Chen, Pin-Yu Chen, Yanzhi Wang, and Xue Lin. Adversarial t-shirt! evading person detectors in a physical world, 2020b. URL <https://arxiv.org/abs/1910.11099>.

Yuan Yao, Tianyu Yu, Ao Zhang, Chongyi Wang, Junbo Cui, Hongji Zhu, Tianchi Cai, Haoyu Li, Weilin Zhao, Zhihui He, et al. Minicpm-v: A gpt-4v level mllm on your phone. *arXiv preprint arXiv:2408.01800*, 2024.



## A Extended Experimental and Implementation Details

To ensure the full reproducibility of our Physical Sharpness-Aware Minimization (P-SAM) framework and to provide deeper insights into the physical deployment constraints, we elaborate on the intricate experimental setups, hardware specifications, and optimization hyperparameters omitted from the main text due to space limitations.

### A.1 Evaluation Details

**(1) Simulation Setting:** Each suite consists of 10 tasks, with each task executed for 50 trials, resulting in a total of 500 rollouts, following [Kim et al., 2024]. For reproducibility, we carefully selected patch paste locations for each suite, ensuring they did not obscure objects or robots in the test environment. To emulate physical degradation digitally, we injected Gaussian sensor noise ( $\sigma = 0.05$ ) and dynamic motion blur to the camera observations at each timestep. **(2) Real-World Setting:** We deployed VLAs on a physical 7-DoF PiPER robotic arm equipped with a Robotiq Hand-E Gripper and a static RGB shoulder camera. The evaluation encompassed 100 trials for each of the three distinct tasks: object grasping, placement, and manipulation.

### A.2 Compute Platform and Hardware Infrastructure

**Offline Optimization Platform.** All digital adversarial optimizations, including the macroscopic Expectation over Transformation (EoT) sampling and the microscopic minimax gradient computations (P-SAM), were conducted on a dedicated computing cluster equipped with NVIDIA RTX 5880 GPUs (48GB VRAM per GPU) to accommodate the intensive memory requirements.

**Software Stack.** The optimization pipeline was implemented on Ubuntu 22.04 LTS using Python 3.10 and PyTorch 2.3.0 with CUDA Toolkit 12.1. During optimization, frozen VLA models such as OpenVLA-7B were loaded via the Hugging Face transformers and accelerate libraries, while the diffusion-based  $\pi_0$  was deployed using its respective custom environment. All model forward passes and gradient computations were executed in bfloat16 precision. This precision choice optimally balances VRAM utilization while strictly preserving gradient fidelity (preventing underflow) during the deep backward pass to the image space. Furthermore, FlashAttention-2 and gradient checkpointing were enabled to significantly accelerate the attention mechanisms within the visual encoders (e.g., SigLIP) and LLM backbones, while strictly containing the activation memory footprint. Finally, spatial physical augmentations were hardware-accelerated using the kornia library, keeping the entire EoT (Expectation over Transformation) pipeline fully differentiable and strictly on the GPU without CPU-transfer bottlenecks.

### A.3 Real-World Robotic Platform Configuration

**Robotic Manipulator.** We utilized the PiPER 7-DoF robotic arm, equipped with a parallel-jaw gripper. This hardware platform closely mimics the kinematic structures and joint compliance seen in common industrial and assistive collaborative robots (cobots).

**Physical Fabrication Constraints.** To ensure strict adherence to the Non-Printability Score (NPS) constraints, all adversarial patches were exported as high-resolution (300 DPI) PDFs and printed using a standard commercial HP Tank 672 printer utilizing standard CMYK ink profiles. To simulate everyday, real-world deployment scenarios, we strictly printed the patches on ordinary A4 office paper. Glossy photo paper was deliberately avoided to minimize unpredictable specular reflections (glare) under direct laboratory lighting, which introduces severe out-of-distribution pixel clipping that is entirely unmodeled by standard continuous noise assumptions.

### A.4 Detailed Optimization Hyperparameters

**Initialization and Optimizer.** The patch parameters  $\delta$  were initialized with neutral gray values (RGB value 0.5 across all channels) to prevent early clipping and maximize the bidirectional dynamic range for pixel updates. We optimized the patch for  $K = 2,500$  iterations using the standard Adam optimizer with a base learning rate of  $\alpha = 0.01$ . The learning rate was decayed using a Cosine Annealing schedule down to  $1 \times 10^{-4}$  to explicitly prevent gradient overshoot and ensure stable convergence within the sculpted flat local minima.

595 **Macroscopic EoT Distribution ( $\Gamma$ ).** For each iteration, we sampled a mini-batch of  $\mathcal{B} = 16$  clean  
 596 frames from the offline datasets. The continuous physical transformation distribution encompassed:  
 597 (1) Scale adjustments uniformly sampled from  $[0.75, 1.25]$  to simulate the arm moving closer to or  
 598 further from the patch; (2) In-plane rotations sampled from  $[-25^\circ, 25^\circ]$  to account for base arm  
 599 orientations; (3) Perspective warping, shifting the four corners of the patch independently by up to  
 600 15% of its spatial dimensions to simulate off-axis 3D viewing angles (yaw/pitch); (4) Brightness  
 601 and contrast multipliers uniformly sampled from  $[0.85, 1.15]$  to emulate ambient shadow casting.

602 **Fabricability Constraint Weights.** The Total Variation (TV) penalty weight was empirically set  
 603 to  $\lambda_1 = 1.5 \times 10^{-4}$ . Excessive TV weight leads to overly simplistic, monochrome patches that  
 604 lack semantic features, while insufficient weight results in high-frequency static that fails to survive  
 605 real-world motion blur. The Non-Printability Score (NPS) weight was set to  $\lambda_2 = 2.0 \times 10^{-3}$ . The  
 606 color gamut set  $C$  for NPS was populated with 10,000 discrete, physically reproducible RGB tuples  
 607 densely sampled from the standard Fogra39 CMYK ICC profile.

## 608 B Extended Related Work and Differential Analysis

609 **Adversarial Attacks in Robotics.** With the widespread application of Vision-Language-Action  
 610 (VLA) models in end-to-end robotic control, their adversarial vulnerabilities have increasingly be-  
 611 come a research focus. Because VLA models directly map cross-modal inputs to low-level physical  
 612 actions, conventional digital-domain attacks often fail as they disregard the robot’s kinematic con-  
 613 straints. To address this, recent studies [Wang et al., 2025a, Lu et al., 2026] have designed phys-  
 614 ical attack objectives tailored to the geometric characteristics and degrees of freedom (DoFs) of  
 615 robotic arms. By maximizing spatial and action discrepancies (e.g., UADA, UPA) or directionally  
 616 manipulating the end-effector trajectory (TMA), these methods successfully expose the execution  
 617 vulnerabilities of VLAs in the physical space.

618 **Inverting Sharpness-Aware Minimization (SAM).** Originally introduced by Foret et al. [2021],  
 619 Sharpness-Aware Minimization (SAM) is a regularization technique designed to improve the gener-  
 620 alization of deep neural networks during standard training. By explicitly penalizing sharp minima  
 621 in the weight space, SAM ensures that model weights converge to flat regions, thereby reducing  
 622 overfitting. A few recent works have attempted to apply SAM to adversarial training to build more  
 623 robust classifiers. Our framework fundamentally inverts and repurposes SAM for offensive physical  
 624 deployment. First, standard SAM perturbs the model weights  $\theta$ . An attacker targeting a deployed  
 625 physical robot has no white-box access to modify the weights of a frozen VLA. P-SAM theoretically  
 626 reformulates the objective to operate strictly within the input perturbation space  $\delta$ . Second, existing  
 627 digital flat-minima techniques apply perturbations homogeneously. In contrast, P-SAM computes a  
 628 conditionally worst-case approximation  $\epsilon^*(x, \tau)$  that is dynamically tied to the macroscopic spatial  
 629 transformation  $\tau$  at that exact iteration. This means P-SAM acts as an ensemble flat-minima seeker;  
 630 it actively sculpts a loss landscape that is universally flat across all possible physical viewpoints the  
 631 moving robot might traverse, rather than just locally flat in a singular static frame.

## 632 C Theoretical Insights into Vulnerability Basin Sculpting

633 To further substantiate the empirical success of P-SAM, we provide a deeper theoretical intuition  
 634 regarding the mechanics of adversarial flickering, the justification for our Taylor approximation, and  
 635 the critical necessity of opaque masking.

### 636 C.1 The Mathematical Mechanics of Adversarial Flickering

637 In standard adversarial patch generation, optimizers seek to minimize the objective  $\mathcal{L}_{base}(\delta)$  with  
 638 respect to the patch pixels  $\delta$ . Mathematically, standard Empirical Risk Minimization (ERM) in-  
 639 evitably converges to a local minimum  $\hat{\delta}$  characterized by a sharp valley. At this minimum, the Hes-  
 640 sian matrix  $H = \nabla_{\delta}^2 \mathcal{L}_{base}(\hat{\delta})$  exhibits extremely large positive eigenvalues. In a physical robotic  
 641 deployment, the true observation is modulated by physical sensor noise and microscopic motion  
 642 blur, denoted as a noise vector  $\eta$ . The effective patch fed to the VLA becomes  $\hat{\delta} + \eta$ . According to  
 643 the second-order Taylor expansion:

$$\mathcal{L}_{base}(\hat{\delta} + \eta) \approx \mathcal{L}_{base}(\hat{\delta}) + \eta^\top \nabla_{\delta} \mathcal{L}_{base}(\hat{\delta}) + \frac{1}{2} \eta^\top H \eta. \quad (9)$$

Since the first-order derivative  $\nabla_{\delta} \mathcal{L}_{base}(\hat{\delta}) \approx 0$  at the local minimum, the loss shift is dominated by the second-order term. Because the eigenvalues of  $H$  are massive at a sharp minimum, even an imperceptible microscopic physical noise vector  $\eta$  causes the adversarial loss to surge quadratically, instantly neutralizing the attack and allowing the robot to self-correct. This mathematically formalizes the phenomenon of adversarial flickering.

## C.2 Justification of the First-Order Taylor Approximation in P-SAM

In our minimax formulation (Equation 3), finding the exact worst-case microscopic noise  $\epsilon$  that maximizes the base loss within an  $\ell_2$  ball of radius  $\rho$  is a highly non-convex inner optimization problem. Solving this via iterative Projected Gradient Descent (PGD) at every single step of the outer optimization would multiply the computational cost by a factor of  $N$ , rendering patch generation prohibitively slow for 7B+ parameter models. Assuming the loss landscape is locally smooth within the highly constrained microscopic radius  $\rho$ , we approximate the perturbed loss using a first-order Taylor expansion:

$$\mathcal{L}_{base}(\delta + \epsilon, x, \tau) \approx \mathcal{L}_{base}(\delta, x, \tau) + \epsilon^{\top} \nabla_{\delta} \mathcal{L}_{base}(\delta, x, \tau). \quad (10)$$

To maximize this linear approximation subject to  $\|\epsilon\|_2 \leq \rho$ , the optimal solution  $\epsilon^*$  strictly aligns with the direction of the local gradient. This elegantly yields our closed-form approximation in Equation 4. Because  $\rho$  represents subtle physical hardware noise rather than global structural semantic shifts, the local linearity assumption holds remarkably well, allowing P-SAM to dynamically track the worst-case degradation with minimal computational overhead.

## C.3 Physical Realizability Constraints of the Dynamic Opaque Mask

A critical innovation in our Macroscopic Kinematic Simulation (Section 2.1) is the mathematical enforcement of the dynamic opaque mask  $M_{\tau} \in \{0, 1\}^{H \times W}$ . Rather than acting merely as a spatial bounding box,  $M_{\tau}$  serves as the mathematical guarantor of physical realizability. To distinguish a legitimate localized physical attack from an impossible global digital perturbation,  $M_{\tau}$  enforces two strict physical constraints:

**(1) Strict Opaque Occlusion:** Standard digital attacks typically utilize an additive perturbation formulation  $x_{adv} = x + \delta$  or alpha-blending  $x_{adv} = (1 - \alpha)x + \alpha\delta$ . In the digital domain, this mathematically models a translucent, “ghost-like” overlay, allowing the optimizer to inadvertently exploit residual background pixels (e.g., the texture of a wooden table or the edges of the target object) underlying the patch to minimize the loss. However, in the physical world, a printed piece of paper is 100% opaque ( $\alpha = 1$ ). This geometric discrepancy destroys the adversarial gradient alignment upon physical printing, causing immediate sim-to-real failure. To resolve this, our objective  $o(x, \delta, \tau) = (1 - M_{\tau}) \odot x + M_{\tau} \odot \mathcal{T}(\delta, \tau)$  acts as a precise spatial carving mechanism. The term  $(1 - M_{\tau}) \odot x$  completely zeroes out the background pixels before patch insertion. This forces the optimizer to achieve targeted task failure utilizing solely the information capacity of the patch pixels themselves, ensuring theoretical fidelity between the simulated optimization state and the physical deployment state.

**(2) Dynamic Geometric Synchronization:** As the robot’s camera translates through 3D space, the 2D projection of the physical patch continuously distorts (e.g., a square sticker viewed from an acute angle projects as a trapezoid or rhombus). By subjecting the base binary mask  $M$  to the exact same macroscopic transformation  $\tau$ , the resulting mask  $M_{\tau}$  dynamically tracks the valid geometric boundary of the perturbation. This strict synchronization prevents the introduction of out-of-bound digital artifacts (such as the zero-padding artifacts inherently introduced by differentiable spatial transformations during affine rotations or perspective warps), ensuring the adversarial gradients are strictly confined to the physical dimensions of the deployable sticker.

## D Supplementary Experiments

### D.1 Resilience Against Algorithmic Input Defenses

A standard, intuitive countermeasure against visual adversarial attacks involves input-level pre-processing. As shown in Figure 9, to evaluate the robustness of our framework against such mitiga-




Figure 9: Task Failure Rates (TFR) under low-pass algorithmic defenses. (a) Our P-SAM maintains high efficacy. (b) Conventional unconstrained baseline suffers a precipitous drop.




Figure 10: Task Failure Rates (TFR) under explicit semantic defense instructions. (a) Our P-SAM efficiently overrides semantic guardrails. (b) Conventional unconstrained baseline.

tions, we subjected the VLA’s visual observations to conventional image transformations in simulation, specifically JPEG compression and Gaussian blurring. Because conventional unconstrained baselines heavily rely on brittle, high-frequency pixel artifacts, these simple low-pass algorithmic defenses easily neutralized the adversarial signals, triggering a precipitous drop in their Task Failure Rate (TFR). In contrast, the P-SAM patch demonstrated inherent resilience to these visual interventions, benefiting from organic and locally uniform macroscopic structures forged jointly by the  $\mathcal{L}_{TV}$  constraint and minimax optimization. This empirical evidence confirms that actively sculpting a flat-minima geometry fundamentally circumvents standard filtering-based robustification techniques.

## D.2 Visual Hijacking vs. Explicit Semantic Instructions

Given that VLAs inherently rely on multimodal vision-language alignment, we investigated whether explicit semantic priors could act as a cognitive defense to override our physical visual attack. As shown in Figure 10, we designed a semantic override scenario where the VLA was commanded via natural language with explicit textual guardrails: “Pick up the target object, and strictly ignore the colored patch on the table.” Strikingly, the P-SAM patch successfully induced a state of “language deafness.” The VLA completely disregarded the explicit semantic defense and compulsively executed the maliciously injected kinematic trajectory. This exposes a profound modality dominance vulnerability in current embodied foundation models: persistent, maliciously optimized visual stimuli can overwhelmingly dominate and bypass language-conditioned safety alignments, dictating physical actions regardless of semantic constraints.

## 711 E Real-World Evaluation Scenarios

712 **Object Grasping (OBJ):** As the most fundamental primitive (e.g., “Pick up the blue cube”), this  
713 task requires precise depth estimation and target localization. The adversarial patch is placed ad-  
714 jacent to the target object. The goal is to corrupt the VLA’s spatial reasoning. Successful attacks  
715 typically manifest as the VLA grasping thin air, closing the gripper prematurely, or forcefully col-  
716 liding with the table.

717 **Object Placement (PLC):** This task involves both picking and complex spatial positioning (e.g.,  
718 “Place the block into the bowl”), demanding robust hand-eye coordination. The patch continuously  
719 perturbs the spatial understanding of the receptacle, frequently forcing the robot to release the object  
720 entirely outside the bowl or deliberately drop it off the edge of the workspace.

721 **Complex Manipulation (MNP):** This task involves articulated objects and continuous environmen-  
722 tal contact (e.g., “Push the sliding door to the left”). Adversarial flickering is most prominent here,  
723 as the physical contact dynamics and mechanical friction cause severe micro-vibrations in the cam-  
724 era mount. Standard static patches fail completely under these vibrations, whereas P-SAM’s flat  
725 basin seamlessly absorbs this hardware jitter, maintaining a steady adversarial drift that causes the  
726 robot to miss the handle entirely.

## 727 F Broader Impacts and Potential Defensive Paradigms

### 728 F.1 Dual-Use Nature and Ethical Disclosure

729 The development of physical-world adversarial attacks against robotic systems inherently carries  
730 dual-use implications. As VLA foundation models accelerate the deployment of autonomous agents  
731 in unstructured human environments (e.g., domestic service robots, automated manufacturing, and  
732 healthcare assistants), ensuring their operational safety becomes a life-critical imperative.

733 This research explicitly highlights a profound, previously underappreciated vulnerability: closed-  
734 loop robotic systems can be completely hijacked, overriding safety protocols and semantic instruc-  
735 tions, simply by passively placing a carefully generated, printable sticker in their visual workspace.  
736 While our P-SAM framework provides a highly effective blueprint for generating physical attacks  
737 without requiring digital access to the robot’s network, this serves as a necessary “red-teaming”  
738 effort. We disclose these vulnerabilities strictly for academic auditing, robustness evaluation, and  
739 defensive research purposes, aiming to catalyze the development of secure robotic systems before  
740 widespread commercial integration occurs.

### 741 F.2 Proposed Defensive Countermeasures

742 **Adversarial Pre-Training in Action Spaces.** Current foundation models are trained exclusively on  
743 pristine trajectory datasets (e.g., Open X-Embodiment). Incorporating P-SAM-generated patches  
744 directly into the massive pre-training loop as a data augmentation technique could force the VLA’s  
745 vision encoder to learn to semantically disentangle high-frequency adversarial noise from valid task-  
746 relevant geometric features. However, scaling this to billions of parameters currently presents a  
747 severe computational bottleneck.

748 **Multimodal Sensory Fusion and Cross-Validation.** Over-reliance on single-modality RGB sen-  
749 sors (a single wrist or third-person camera) creates a critical single point of failure. Integrating a  
750 multi-camera fusion setup could allow the robotic system to cross-verify the 3D geometry of the  
751 scene. An adversarial patch is highly unlikely to perfectly satisfy the worst-case flat-minima op-  
752 timization for two distinctly separated physical viewpoints simultaneously. Furthermore, fusing  
753 visual inputs with active depth sensing (LiDAR) or tactile feedback would drastically increase the  
754 difficulty for an attacker to generate a viable physical override.

## 755 G More Visual Presentations

756 **More attack examples.** Figure 11 shows more successful results of our attacks, while Figure 12  
757 presents a failed result. The model we have deployed here is OpenVLA [Kim et al., 2024]. As il-  
758 lustrated in Section 5, severe topological occlusions of the camera’s line-of-sight emerge as a highly






Figure 11: Examples of successful physical attacks across different manipulation tasks.




Figure 12: An example of a failed attack rollout. The task objective is to put the apple on the plate. We can observe that when the robotic arm accidentally covers the adversarial patch with the plate, the transmission of adversarial visual features is severed, causing the attack to fail.

759 probable cause of physical attack failures. Because closed-loop systems rely on continuous visual  
 760 streaming to update action predictions, the adversarial influence of the P-SAM patch demands an  
 761 uninterrupted visual pathway. Momentary loss of visibility abruptly halts the injection of worst-case  
 762 pixel perturbations into the vision encoder, rendering the sculpted flat vulnerability basins inac-  
 763 cessible. Consequently, the adversarial manipulation sequence is suspended, allowing the model  
 764 to momentarily recover its benign control policy and break the hijacking loop. This phenomenon  
 765 confirms that while P-SAM successfully mitigates algorithmic adversarial flickering, continuous ge-  
 766 ometric visibility remains a strict, physical prerequisite for reliably executing the attack in the real  
 767 world.

768 **More patch examples.** Figure 13 visualizes a selection of our generated adversarial patches. As  
 769 can be observed, our patches exhibit distinct, organic macroscopic structures. This unique visual  
 770 appearance is a direct consequence of our optimization framework: the Total Variation ( $\mathcal{L}_{TV}$ ) con-  
 771 straint, coupled with the microscopic minimax optimization (P-SAM), jointly forge these piecewise-




Figure 13: Examples of generated P-SAM adversarial patches. The target model is OpenVLA.

772 constant macro-patterns. By sculpting these flat vulnerability basins, our patches inherently resist  
 773 common image-space interventions in the physical world (such as lighting variations or viewpoint  
 774 deformations), ensuring robust deployability and high transferability across different VLA architec-  
 775 tures.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: Sec. 2, 3, 4

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Sec. 5

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: Sec. 2 and App. C

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Sec. 3 and App. A

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

#### 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: See our project website

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Sec. 3 and App. A

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Sec. 3 and App. D

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.



- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: App. A

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: All sections

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Sec. 1, 4, 5 and App. F

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [\[Yes\]](#)

Justification: App. F

Guidelines:

- The answer [\[N/A\]](#) means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: Sec. 3

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

### 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: See our project website

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: [\[N/A\]](#)

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: [\[N/A\]](#)

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

### 16. Declaration of LLM usage

1084 Question: Does the paper describe the usage of LLMs if it is an important, original, or  
1085 non-standard component of the core methods in this research? Note that if the LLM is used  
1086 only for writing, editing, or formatting purposes and does *not* impact the core methodology,  
1087 scientific rigor, or originality of the research, declaration is not required.

1088 Answer: [N/A]

1089 Justification: [N/A]

1090 Guidelines:

- 1091 • The answer [N/A] means that the core method development in this research does not  
1092 involve LLMs as any important, original, or non-standard components.
- 1093 • Please refer to our LLM policy in the NeurIPS handbook for what should or should  
1094 not be described.

1095